

Priyanshi Sharma

Jaipur, Rajasthan | priyanshi.sharma1974@gmail.com | +91 9928019445
github.com/PRI4294 | linkedin.com/in/priyanshi333 | leetcode.com/u/_priyanshi_

Profile Summary

B.Tech in Information Technology with production-level AI engineering experience across LLM orchestration, RAG pipelines, real-time voice systems, and Text-to-SQL engines. Has built and deployed AI products in enterprise and startup environments, achieving measurable outcomes including 80%+ query accuracy and zero-downtime voice pipelines. Proficient in Python, PyTorch, Generative AI, HuggingFace Transformers, LangChain, and FastAPI. Focused on shipping AI systems that are reliable and production-ready.

Experience

Biz4Group LLC

Jaipur | Feb 2026 – Present

AI Engineer

- Designed and shipped production AI systems integrating LLMs, voice APIs, and third-party REST services — automating complex multi-step workflows end-to-end and reducing manual operational effort to zero across deployed products.
- Engineered deterministic LLM tool-calling orchestration with schema-validated API payloads, resolving 12+ production failure modes (400/401/404) and reducing system error rate from ~100% to 0% through strict input validation, guardrails, and fault-tolerant pipeline design.
- Built and deployed real-time AI voice systems using expressive LLM-driven voice models, low-latency WebRTC audio streaming, and cloud infrastructure (AWS EC2, S3, PostgreSQL) — optimizing conversational flow, response latency, and session persistence at scale.
- Delivered a Python-based data processing backend handling 5,000+ records across 8 report types, achieving 95% test coverage (65+ pytest cases) and shipping a secure, production-ready executable with input sanitization and path traversal protection.

Hewlett Packard Enterprise

Remote | Feb 2025 – Jul 2025

Project Intern – HPE CPP Program

- Engineered a Text-to-SQL system using GPTQ-quantized Qwen2.5-7B, achieving 80%+ accuracy on 1,500+ real-world enterprise SQL queries and reducing manual data retrieval effort by 60%.
- Designed prompt engineering strategies (chain-of-thought, few-shot) that cut SQL generation errors by 25%, and optimized inference to run within a 5GB memory budget enabling fully offline, local deployment.
- Built a real-time SQL streaming interface with FastAPI and Streamlit, reducing average query wait time by 30% and enabling sub-second response streaming.
- Developed an automated evaluation pipeline using semantic equivalence checks across 600+ test cases, with SQL injection prevention and destructive-command filtering ensuring 100% safe query execution.

Education

B.Tech (Information Technology)

2022 – 2026

JECRC Foundation, Jaipur, Rajasthan | CGPA: 8.74 / 10

Projects

Local Voice AI Agent with RAG

Feb 2026

AI/ML Engineer – Personal Project

- Architected a fully offline, zero-cloud-cost voice AI system — a 3-stage pipeline (STT → LLM → TTS) using Moonshine, Llama 3.2-1B, and Kokoro-82M — achieving ~800ms end-to-end latency with GPU-accelerated inference (PyTorch, CUDA 12.1+) on consumer hardware.
- Integrated a FAISS-based RAG system using HuggingFace Transformers for context-aware, document-grounded responses, with chunked embedding indexing for efficient retrieval — fully local deployment with zero cloud cost and no external API dependencies.
- Implemented GPU VRAM-aware memory management to sustain multi-turn conversations without session degradation, with WebRTC-based browser streaming for real-time audio delivery.

AI Voice Interview Simulator

March 2026

Full-Stack AI Engineer – Personal Project

- Built a cross-platform AI mock interviewer (iOS/Android) that generates role-specific, JD-grounded interview questions in real time — powered by Mistral-7B, prompt chaining via Hugging Face Inference API, and a React Native + FastAPI full-stack architecture.
- Designed a semantic answer evaluation pipeline using HuggingFace sentence embeddings and cosine similarity, outperforming keyword matching by 40%+ on semantic answer evaluation accuracy — enabling nuanced, context-aware scoring beyond surface-level word overlap.
- Stored session data in PostgreSQL to track performance across sessions and surface weak areas per role type, enabling personalized feedback loops for targeted interview preparation.

Skills

Languages & API's: Python, SQL, C++, REST APIs, FastAPI, WebRTC, Git

LLM & GenAI: Generative AI, LLMs, LangChain, LangGraph, RAG, KAG, CAG, Prompt Engineering (CoT, Few-shot), Tool-Calling Orchestration, HuggingFace Transformers, Vector Databases, Embeddings, Similarity Search, OpenAI API

ML & Deep Learning: Machine Learning, PyTorch, TensorFlow, Scikit-learn, Model Training, Fine-tuning (LoRA, PEFT), Model Quantization (GPTQ), Computer Vision, CUDA, Pandas, NumPy, NLTK, spaCy

Backend & Infrastructure: MLOps, Docker, AWS (EC2, S3), PostgreSQL, SQLite, MongoDB, SQLAlchemy, Firebase, Supabase, Streamlit, React Native, Jupyter

CS Fundamentals: DSA, OOPs, DBMS, OS, Computer Networks